

1 Linear Tests

1.1 Joint Significance of Subset of Regressors

Suppose we run the regression

$$y = \beta_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + u,$$

and slope coefficients for both x_3 and x_4 have high enough p -values that we conclude each one is statistically insignificant. It is still possible, however, that they are *jointly* significant, even if they are individually insignificant. In other words, we want to test

$$H_0 : \beta_3 = \beta_4 = 0,$$

$$H_1 : \text{at least one of } \beta_3, \beta_4 \neq 0.$$

Think of H_0 as being a *restriction* placed on β_3 and β_4 that we want to test.

The first thing to do is take the model where β_3 and β_4 are unrestricted (that is, a regression where x_3 and x_4 are included and thus their coefficients are estimated) and find its sum of squared residuals, call it RSS_{ur} . Then make the restrictions (by not even including them in the regression, which implicitly sets the slope coefficients equal to zero) and find the sum of squared residuals for this restricted model, call it RSS_r .

If β_3 and β_4 are jointly insignificant, i.e. if H_0 is true, then you would expect the difference between the two RSS terms to be small since the RSS represents unexplained variation in y . In other words, if β_3 and β_4 are jointly insignificant, then we shouldn't expect much difference in how well the model explain things whether they're both simultaneously included or not. Also notice that RSS_{ur} is *always* smaller than RSS_r , warranting a one-sided test: we reject the null if the difference between RSS_r and RSS_{ur} is really big, which implies that the restricted model leaves a lot unexplained.

We just need to formalize what we mean by a “big” difference between the two. This is given by the F -statistic

$$F \equiv \frac{(RSS_r - RSS_{ur})/(q)}{RSS_{ur}/(n - k)} \sim F_{q, n-k},$$

where

- n is the number of observations;
- k is the number of parameters being estimated in the unrestricted model, in this case

$k = 4$ because we estimate the intercept plus slope coefficients for x_2 , x_3 , and x_4 ;

- q is the number of parameters being tested (i.e. restrictions), in this case $q = 2$;
- $F_{q,n-k}$ is the F -distribution with q parameters restricted parameters and $n - k$ is the unrestricted degrees of freedom.

Sometimes g is used to denote the number of estimates being made in the restricted regression, and hence $q = k - g$. But I find that confusing so I'm going to stick with q .

1.2 Overall Significance of Regressors

At the extreme end, we can also test whether *all* regressors are jointly significant by comparing it to a regression with *no* regressors. There are k things being estimated in the full regression, and we still estimate the intercept in the restricted regression, so we are restricting $q = k - 1$ parameters. This yields the F -statistic

$$F \equiv \frac{R^2/(k-1)}{(1-R^2)/(n-k)} \sim F_{k-1,n-k},$$

where R^2 is given from the unrestricted regression. This is the F -statistic given in `stargazer` output.

1.3 Individual Significance of Regressor

At the other extreme end, we can use the F -test to test just one restriction, e.g. $H_0 : \beta_3 = 0$ against $H_1 : \beta_3 \neq 0$. This looks like a simple, ordinary hypothesis test, and indeed, the F -statistic in this case will be the usual t -statistic squared because $F(1, n-k) = [T(n-k)]^2$.

This joint significance test, including both extreme cases, comes with an important caveat: we must have homoskedasticity in order for the test to be valid.

1.4 Example

Go to my website and load `wages.csv` into R. We run the regression

$$wage = \beta_1 + \beta_2 educ + \beta_3 IQ + \beta_4 sibs + \beta_5 brthord + u,$$

where *sibs* is the number of siblings a person has and *brthord* is that person's order of birth, e.g. *brthord* = 1 if the person is the first-born child of the family.

TABLE 1

<i>Dependent variable:</i>	
	wage
educ	42.095*** (6.909)
IQ	4.650*** (1.030)
sibs	-6.395 (7.118)
brthord	-11.750 (10.057)
Constant	-26.029 (108.522)
Observations	852
R ²	0.142
Adjusted R ²	0.138
F Statistic	34.950*** (df = 4; 847)

Note: *p<0.1; **p<0.05; ***p<0.01

As a matter of accounting, note that $n = 852$ and we have $k = 5$ estimates in the unrestricted model: four variables plus the constant.

We can see from the absence of asterisks that neither sibling count nor birth order are statistically significantly different from zero. However we might still suspect that the two variables are jointly significant. To quell our suspicion, we run the restricted regression

$$wage = \beta_1 + \beta_2 educ + \beta_3 IQ + u$$

apropos the joint significance test

$$H_0 : \beta_4 = \beta_5 = 0,$$

$$H_1 : \text{at least one of } \beta_4, \beta_5 \neq 0.$$

Again as a matter of accounting, we are making $q = 2$ restrictions. So we know we will be using the

$$F(q, n - k) = F(2, 852 - 5)$$

distribution to judge the test statistic.

So uh, let's find the test statistic. R allows us to see the residuals of each observation, so we must square and then sum them to get RSS . The code is below, but we find

$$RSS_{ur} = 119,125,941,$$

$$RSS_r = 119,827,678.$$

Now let's plug this into the formula for the F -statistic, which yields

$$F = \frac{(119,827,678 - 119,125,941)/(2)}{119,125,941/(847)} = 2.495.$$

This gives a p -value of `pf(2.495, 2, 847, lower.tail=FALSE) = 0.083`. Hence we fail to reject the null at 5% level; our F -statistic isn't big enough to reject. Thus we conclude that variables *sibs* and *brthord* are individually *and* jointly insignificant at 5% significance level. At the 10% level, however, we can reject the null: while *sibs* and *brthord* are both individually insignificant, they are jointly significant at 10% significance level.

Alternatively, we can use command `linearHypothesis()` from the `car` package in R.

```

1 library("stargazer")
2 library("car")
3
4 wages <- read.csv("wages.csv")
5
6 regur <- lm(wage ~ educ + IQ + sibs + brthord, data = wages)
7 regr  <- lm(wage ~ educ + IQ, data = wages)
8
9 RSSur = sum(regur$residuals^2)          ### unrestricted RSS
10 RSSr  = sum(regr$residuals^2)          ### restricted RSS
11
12 F     = ((RSSr - RSSur)/2) / (RSSur/(847))  ### F-statistic
13 pv    = pf(F, 2, 847, lower.tail=FALSE)    ### p-value
14
15 Hnull <- c("sibs=0", "brthord=0")          ### null hypothesis
16 linearHypothesis(regur, Hnull)             ### joint significance test

```

1.5 Other Linear Restrictions

We can use the same idea to test more complex linear restrictions. We might want to test $H_0 : \beta_2 + \beta_4 + \beta_{15} + \beta_{99} = 4$ for some reason. These can be tedious with pencil and paper, but are easy to do in R by writing the `Hnull` vector as `c("x2 + x4 + x15 + x99 = 4")`.

The way to solve it by pencil and paper is to solve the restriction for one β and then substitute it out of the regression entirely. For instance, consider the model

$$y = \beta_1 + \beta_2 x_2 + \beta_3 x_3 + u.$$

If our restriction is $H_0 : \beta_2 + \beta_3 = 7$, then we can write it as $\beta_3 = 7 - \beta_2$. We can now substitute this in for β_3 and rewrite the model such that

$$\begin{aligned} y &= \beta_1 + \beta_2 x_2 + (7 - \beta_2)x_3 + u \\ \implies y - 7x_3 &= \beta_1 + \beta_2(x_2 - x_3) + u. \end{aligned}$$

You get RSS_r from regressing $y - 7x_3$ on $x_2 - x_3$, both of which you'll probably have to define in R. Then proceed with the F -statistic as given above with $q = 1$ restriction.

2 Breusch-Pagan Test for Heteroskedasticity

Homoskedasticity is a pretty strong condition, so it is definitely something we should check for before we just go around assuming it. To that end, we will use the **Breusch-Pagan test**, described as follows.

2.1 Theory

We will maintain OLS assumptions 1-2 so that estimates are unbiased and consistent, and also note that the Breusch-Pagan test requires normality of disturbances. Our null hypothesis is that of homoskedasticity, written explicitly as

$$\begin{aligned} H_0 : \text{Var}(u_i | x_{2i}, \dots, x_{ki}) &= \sigma_u^2 \quad \text{for all } i, \\ H_1 : \text{Var}(u_i | x_{2i}, \dots, x_{ki}) &\neq \sigma_u^2 \quad \text{for some } i. \end{aligned}$$

Take my word for it that we can express this variance as

$$\text{Var}(u_i | x_{2i}, \dots, x_{ki}) = E[u_i^2 | x_{2i}, \dots, x_{ki}] - E[u_i | x_{2i}, \dots, x_{ki}]^2,$$

which is useful because $E[u_i | x_{2i}, \dots, x_{ki}] = 0$ by zero conditional mean assumption,¹ and furthermore because our regressions have the same conditional expectation interpretation.

¹It is a general formula that $\text{Var}(x) = E[x^2] - E[x]^2$. It is not a difficult result to prove. Try it as an exercise if you are bored enough. I suspect you are not.

Which is to say, we can instead consider

$$H_0 : E[u_i^2 | x_{2i}, \dots, x_{ki}] = \sigma_u^2 \quad \text{for all } i,$$

$$H_1 : E[u_i^2 | x_{2i}, \dots, x_{ki}] \neq \sigma_u^2 \quad \text{for some } i,$$

where regressing u^2 on x_2, \dots, x_k gives us precisely $\hat{u}_i^2 = E[u_i^2 | x_{2i}, \dots, x_{ki}]$. If this regression has overall significance, it means the regressors explain u^2 , and hence $\text{Var}(u^2)$ depends on x_1, \dots, x_k ; in other words, σ_u^2 is not constant, so it is not homoskedastic.

2.2 Algorithm

Without further ado, here is the algorithm for the test.

Step 1. Estimate your model,

$$y = \beta_1 + \beta_2 x_2 + \dots + \beta_k x_k + u,$$

using the typical OLS rigmarole.

Step 2. Calculate the squared residuals \hat{u}_i^2 for each i .

Step 3. Regress \hat{u}^2 on each regressor,

$$\hat{u}^2 = \delta_1 + \delta_2 x_2 + \dots + \delta_k x_k + v,$$

and make note of the R -squared of this regression, call it $R_{\hat{u}^2}^2$. This is called an **auxiliary regression** because we only do it to help analyze the primary regression of step 1.

Step 4. Calculate the F -statistic for the overall significance of the auxiliary regression,

$$F \equiv \frac{R_{\hat{u}^2}^2 / (k - 1)}{(1 - R_{\hat{u}^2}^2) / (n - k)} \sim F_{k-1, n-k},$$

from which you calculate the p -value.

Step 5. Compare the p -value to your chosen level. If the p -value is smaller than your level, then we conclude that some combination of the δ_j parameters have significant effect

on \hat{u}_i^2 . In other words, \hat{u}^2 depends on the values of x_2, \dots, x_k , and therefore we infer that $\text{Var}(u_i|x_2, \dots, x_k)$ does as well, so we can reject homoskedasticity.

2.3 Example

Load up `wages.csv` again. We want to see if the regression

$$\text{wage} = \beta_1 + \beta_2 \text{educ} + \beta_3 \text{IQ} + \beta_4 \text{sibs} + \beta_5 \text{brthord} + u$$

has homoskedastic disturbances or not. Doing some work in R gives us $R_{\hat{u}^2}^2 = 0.0211$. There are $n = 852$ observations and $k = 5$ estimates. Hence we look at the F -statistic

$$F = \frac{0.0211/(4)}{(1 - 0.0211)/(847)} \approx 4.5601.$$

This gives p -value of `pf(4.5601, 4, 847, lower.tail=FALSE) = 0.0012`. Hence we conclude at conventional levels that the disturbance depends on the regressors, and hence we reject homoskedasticity. You could also look at `summary(auxreg)` output to see the p -value.²

```

1 library("stargazer")
2 library("lmtest")
3
4 wages <- read.csv("wages.csv")
5
6 reg <- lm(wage ~ educ + IQ + sibs + brthord, data = wages)
7
8 ### square residuals
9 uhatsq = reg$residuals^2
10
11 ### regress squared residuals
12 auxreg <- lm(uhatsq ~ educ + IQ + sibs + brthord, data = wages)
13
14 ### calculate R squared for auxiliary uhat^2 regression
15 RSSuhat = sum(auxreg$residuals^2)
16 TSSuhat = sum((uhatsq - mean(uhatsq))^2)
17 Rsquhat = 1 - RSSuhat/TSSuhat
18
19 ### calculate F-statistic and p-value
20 F = (Rsquhat/(4)) / ((1 - Rsquhat)/(847))
21 pv = pf(F, 4, 847, lower.tail=FALSE)
22
23 ### be lazy and let R give you the p-value
24 summary(auxreg)
```

²We can use `bptest()` from the `lmtest` package, but it uses a χ^2 test instead of an F -test. For $n - k$ is sufficiently large, it is approximately true that $\chi^2(q)/q = F(q, \infty)$.

3 RESET Test

3.1 Theory

The **RESET test** is used to test whether your model is misspecified or not. The logic is as follows. Suppose we have correctly specified the model using only linear variables. Hence the zero conditional mean condition is satisfied, i.e. $E[u|x_2, \dots, x_k] = 0$, and we can conclude that there are no relevant omitted variables. In particular, we have not omitted any relevant nonlinear functions of the regressors, for instance x_4^2 . Hence if we add some nonlinear aspect to the model, then the corresponding coefficients should be statistically indistinguishable from zero. If not, then we're using the wrong model.

Let's be more explicit. We originally use the model

$$y = \beta_1 + \beta_2 x_2 + \dots + \beta_k x_k + u. \quad (1)$$

We do OLS in the typical fashion and generate fitted values \hat{y} . It is important at this juncture to recognize that \hat{y} is just a function of x_2, \dots, x_k . Accordingly, \hat{y}^2 and \hat{y}^3 and so forth are just nonlinear function of x_2, \dots, x_k . (Squared and cubed terms are most common and useful, so I will stop at the third power.)

The takeaway is that by putting \hat{y}^2 and \hat{y}^3 into the regression, we're including a host of nonlinear terms. Doing so means running regression of form

$$y = \beta_1 + \beta_2 x_2 + \dots + \beta_k x_k + \delta_1 \hat{y}^2 + \delta_2 \hat{y}^3 + v. \quad (2)$$

If the nonlinearities don't matter, then we expect δ_1 and δ_2 to be statistically indistinguishable from zero, and can conclude that our model is not complete garbage. The specific test is of the form

$$H_0 : \delta_1 = \delta_2 = 0,$$

$$H_1 : \text{at least one of } \delta_1, \delta_2 \neq 0.$$

So it's a test of joint significance. We've seen this before. The overall regression equation (2) has $k + 2$ variables (including the intercept). The restricted regression is simply the original regression with k variables because we make two $q = 2$ restrictions in H_0 . Hence we use test statistic

$$F \equiv \frac{(RSS_r - RSS_{ur})/(2)}{RSS_{ur}/(n - k - 2)} \sim F(2, n - k - 2).$$

If the F -statistic is big enough, when we conclude that the model is misspecified in some way because it's saying something important is in \hat{y}^2 or \hat{y}^3 . But we don't know what the important thing is: the big downside to the RESET test is that it doesn't tell us how to proceed when it tells us that our model is junk.

3.2 Example

Load up `wages.csv` again. Let's see if the model

$$wage = \beta_1 + \beta_2 educ + \beta_3 IQ + \beta_4 sibs + \beta_5 brthord + u$$

is misspecified. We run OLS on the preceding equation, we generate variables \hat{y}^2 and \hat{y}^3 , and we throw them into the regression

$$wage = \beta_1 + \beta_2 educ + \beta_3 IQ + \beta_4 sibs + \beta_5 brthord + \delta_1 \hat{y}^2 + \delta_2 \hat{y}^3 + v.$$

Then we test whether δ_1 and δ_2 are jointly significant or not.

There are $n = 852$ observations. In the big model there are $k + 2 = 7$ estimates being made. In the restricted version, there are $k = 5$ estimates being made and $q = 2$ restrictions. Hence we find

$$F \equiv \frac{(RSS_r - RSS_{ur})/(2)}{RSS_{ur}/(852 - 7)} \sim F(2, 845).$$

I'm not going to grind out each RSS , but suffice it to say that R gives $F = 0.6352$. This gives a p -value of `pf(0.6352, 2, 845, lower.tail=FALSE) = 0.5301`. Hence we cannot reject the null at conventional levels and thus we have insufficient evidence of model misspecification. Hooray, the model isn't total garbage!

We can also do this by using the `resettest()` function from the `lmtest` package. By default it will also do second and third powers of \hat{y} . The R code is on the final page.

```
1 library("stargazer")
2 library("car")
3 library("lmtest")
4
5 wages <- read.csv("wages.csv")
6
7 ### run original regression
8 reg <- lm(wage ~ educ + IQ + sibs + brthord, data = wages)
9
10 wages$yhatsq = reg$fitted.values^2      ### generate fitted squared
11 wages$yhatcu = reg$fitted.values^3      ### generate fitted cubed
12
13 ### run RESET regression
14 RESETreg <- lm(wage ~ educ + IQ + sibs + brthord + yhatsq + yhatcu,
15               data = wages)
16
17 RESETRSSur = sum(RESETreg$residuals^2)  ### unrestricted RESET RSS
18 RESETRSSr  = sum(reg$residuals^2)      ### restricted RESET RSS
19
20 ### calculate F statistic and p-value
21 F = ((RESETRSSr - RESETRSSur)/2) / (RESETRSSur/845)
22 pv = pf(F, 2, 845, lower.tail=FALSE)
23
24 ### let R do the restriction testing for you
25 Hnull <- c("yhatsq=0", "yhatcu=0")      ### null hypothesis
26 linearHypothesis(RESETreg, Hnull)       ### joint significance test
27
28 ### let R do the whole damn thing for you
29 resettest(reg)
```